

AUTOJUDGE: PROJECT REPORT

Aayushi Bhardwaj – 23112001

● Project Overview & Problem Statement

Online coding platforms like Codeforces, Kattis, and CodeChef classify programming problems as Easy, Medium, or Hard and assign numerical difficulty scores. However, this process relies on slow human judgment and user feedback.

This project develops an intelligent machine learning system that automatically predicts:

- Problem Class (Easy/Medium/Hard) using classification, and
- Problem Score (numerical difficulty value) using regression,

using only the textual content of the problem, including the problem description, input description, and output description.

The system was trained on a dataset of 4,111 Kattis problems and leverages TF-IDF text features for prediction.

Predictions are delivered through a Streamlit web application, where users can paste new problem text and instantly see the predicted difficulty class and score.

Best Results:

Classification: 52% accuracy using Random Forest

Regression: MAE = 0.743778 , RMSE = 0.920731

● Dataset Description:

Source: https://raw.githubusercontent.com/AREEG94FAHAD/TaskComplexityEvaluation-24/main/problems_data.jsonl

Total problems: 4,111 from Open Kattis

Columns Used:

- title – Problem name
- description – Main problem statement
- input_description – Input format
- output_description – Output format
- problem_class – Easy / Medium / Hard (classification target)
- problem_score – 1.1 to 9.7 (regression target)

Class Distribution:

- Hard: 47% (1,941 problems)
- Medium: 34% (1,405 problems)
- Easy: 19% (766 problems)

df.sample(5)

✓ 0.0s

Python

		title	description	input_description	output_description	sample_io	problem_class	problem_score	url
3938	Symmetric Order	In your job at Albatross Circus Management (ye...	The input consists of one or more sets of stri...	For each input set print "SET\n\$ \$" on a line,...	[{"input": "7 Bo Pat Jean Kevin Claude William..."}]	easy	1.6	https://open.kattis.com/problems/symmetricorder	
3033	ASCII Figure Rotation	Consider an ASCII art figure like the one on t...	Input consists of up to 100 figures. Each figu...	For each figure, print a copy of the figure ro...	[{"input": "5 + + - + + - + +---+ 7..."}]	medium	3.4	https://open.kattis.com/problems/ascifigurerot	
998	Pedal Power	You're ready to start your semester off right!...	The problem input describes a multigraph with ...	Output the minimum possible cumulative travel ...	[{"input": "4 4 0 1 2 3 1 10 2 3 2 2 0 10 4 1 ..."}]	hard	6.9	https://open.kattis.com/problems/pedalpower	
3430	Alien Numbers	The decimal numeral system is composed of ten ...	The first line of input gives the number of ca...	For each test case, output one line containing...	[{"input": "4 9 0123456789 of8 Foo of8 0123456..."}]	easy	2.6	https://open.kattis.com/problems/aliennumbers	
1567	Burrows-Wheeler	The Burrows-Wheeler transform is a technique f...	Input consists of up to \$100\$ messages, one pe...	For each input message, output a single line t...	[{"input": "arbitrary string this is the thing..."}]	hard	6.0	https://open.kattis.com/problems/burrowswheel	

df.shape

✓ 0.0s

(4112, 8)

Python

df.info()

✓ 0.0s

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4112 entries, 0 to 4111
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   title        4112 non-null   object 
 1   description   4112 non-null   object 
 2   input_description  4112 non-null   object 
 3   output_description 4112 non-null   object 
 4   sample_io     4112 non-null   object 
 5   problem_class 4112 non-null   object 
 6   problem_score 4112 non-null   float64 
 7   url          4112 non-null   object 
dtypes: float64(1), object(7)
memory usage: 257.1+ KB
```

• Data Preprocessing

- Missing Values: 120 rows had empty strings ("") in input_description and 131 in output_description, and 81 in description while title had no missing values. These are not true nulls.
- Duplicates: 1 duplicate row was found and removed with df.drop_duplicates(), resulting in a clean dataset of 4,110 rows.

```

#Actual NULL values
print("input description -", (df["input_description"].str.strip() == "").sum())
print("output description -", (df["output_description"].str.strip() == "").sum())
print("description -", (df["description"].str.strip() == "").sum())
print("title -", (df["title"].str.strip() == "").sum())

```

✓ 0s

```

input description - 120
output description - 131
description - 81
title - 0

```

#duplicates

```

df[df.duplicated(subset=["title", "url"], keep=False)]

```

✓ 0s

	title	description	input_description	output_description	sample_io	problem_class	problem_score	url
1233	Advanced Causal Measurements	Causality is a very important concept in theor...	The first line of input is the number of cases...	Output consists of a single line for each case...	[Input]:4 4 1 1 - 1 1 3 1 4 2 6 4 2 1 - 1 - ...	hard	6.5	https://open.kattis.com/problems/causal
4111	Advanced Causal Measurements	Causality is a very important concept in theor...	The first line of input is the number of cases...	Output consists of a single line for each case...	[Input]:4 4 1 1 - 1 1 3 1 4 2 6 4 2 1 - 1 - ...	hard	6.5	https://open.kattis.com/problems/causal

drop duplicates

```

print(df.shape)
df = df.drop_duplicates(subset=["title", "url"])
print(df.shape)

```

✓ 0s

3. Class Imbalance: Distribution of problem classes—Hard: 1,941 (47%), Medium: 1,405 (34%), Easy: 766 (19%)—shows the dataset is highly imbalanced, with hard problems roughly 3x more frequent than easy problems.
4. The dataset has an average difficulty of 5.11/10, with scores ranging from 1.1 (very easy) to 9.7 (very hard) and most problems falling in the medium difficulty range.

```

df.describe()

```

✓ 0s

	problem_score
count	4111.000000
mean	5.114352
std	2.177928
min	1.100000
25%	3.300000
50%	5.200000
75%	6.900000
max	9.700000

df["problem_class"].value_counts()

✓ 0s

problem_class	count
hard	1940
medium	1405
easy	766

Name: count, dtype: int64

df["problem_class"].value_counts()/len(df["problem_class"])

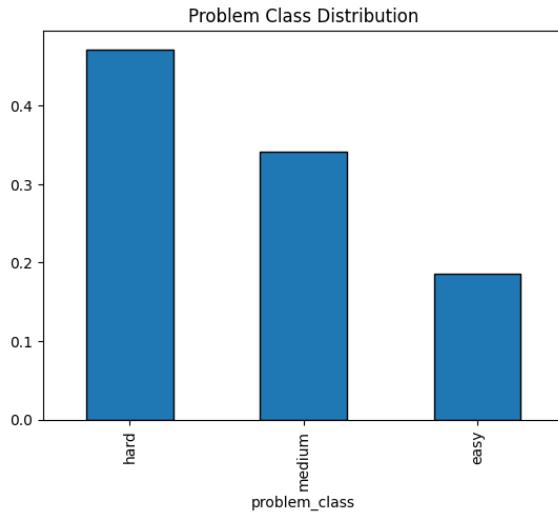
✓ 0s

problem_class	count
hard	0.471905
medium	0.341766
easy	0.186329

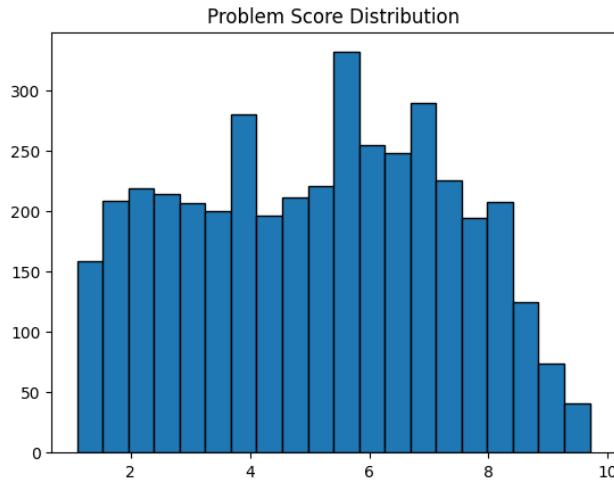
Name: count, dtype: float64

● Data Visualization

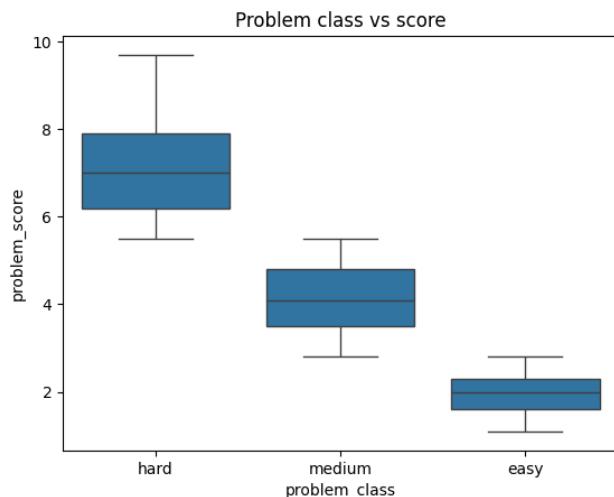
1. Hard problems dominate (47%) while easy problems underrepresented (19%) - highly imbalanced dataset.



2. Most scores cluster around 5 (medium difficulty) with long tail toward harder problems.



3. Hard problems have significantly longer descriptions than easy problems - text length is a strong indicator.



- Feature Engineering:

1. Combined all text

I merged the title, description, input, and output into one single text column called text so the model could see the complete problem statement at once.

2. Text length feature

I calculated the number of characters in each problem.

Easy problems are usually short (~2,000 characters)

Hard problems are longer (~4,000 characters) with more detailed explanations

3. Keyword-based features

I checked whether certain important programming keywords appeared in the text, such as recursion, graph, dp, tree, etc.

For each keyword, I created a new column:

1 → keyword present

0 → keyword absent

4. TF-IDF text features (main features)

I converted the full text into TF-IDF features, resulting in 8,000 numerical features per problem.

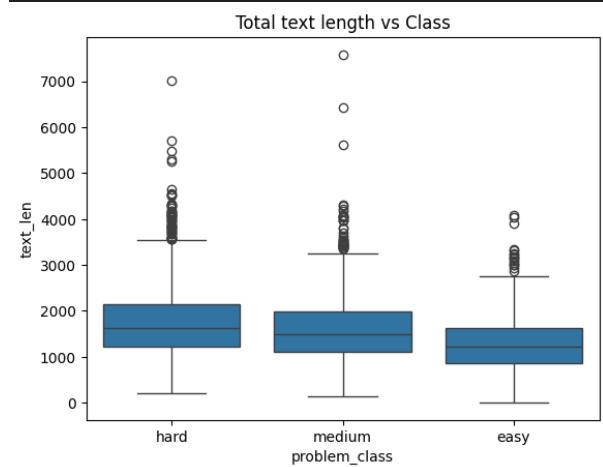
This helps the model understand which words are most important when deciding problem difficulty.

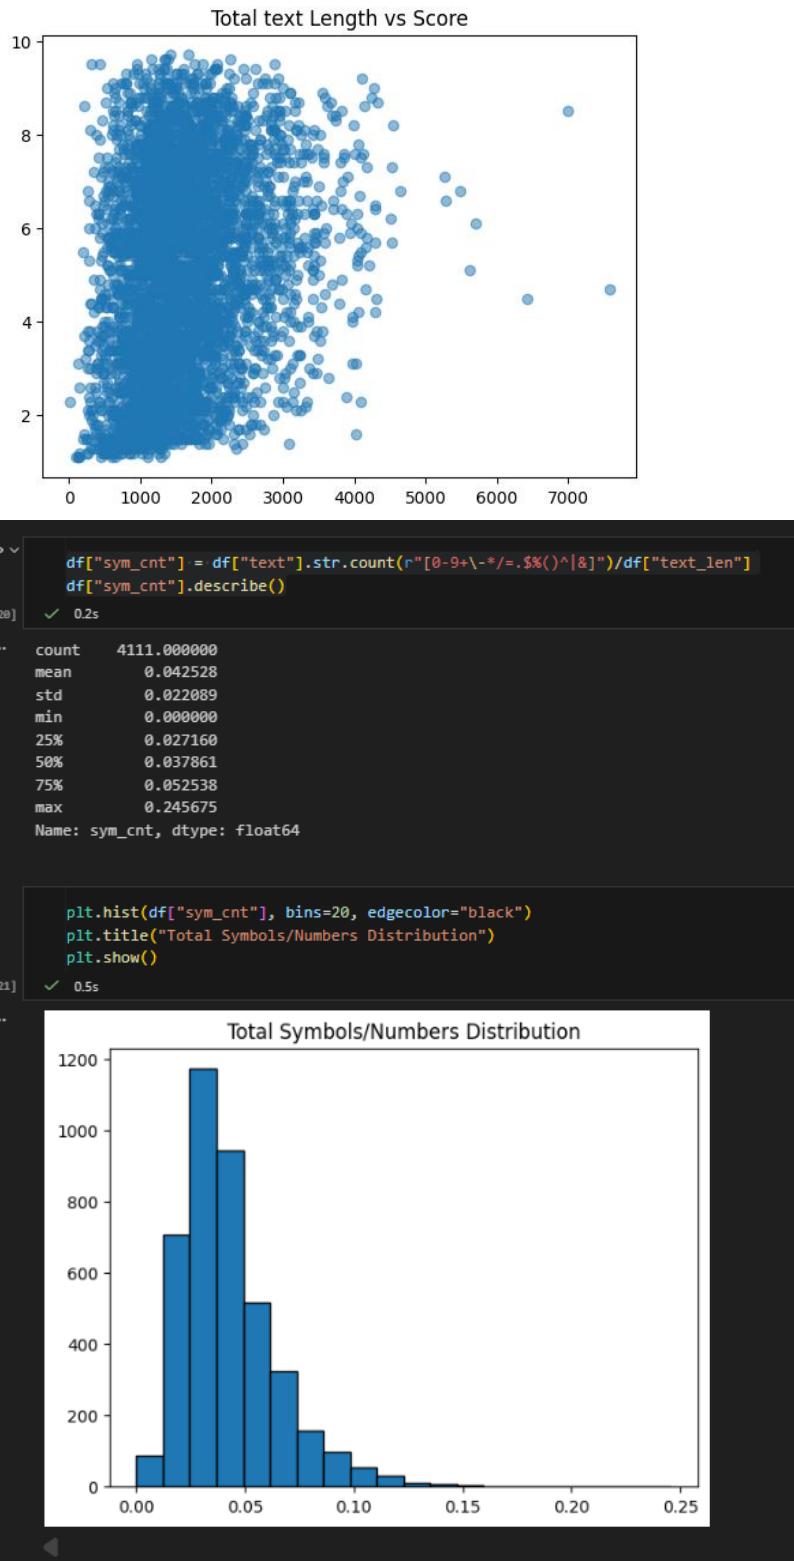
5. I also added variables for word count and unique words count for better understanding

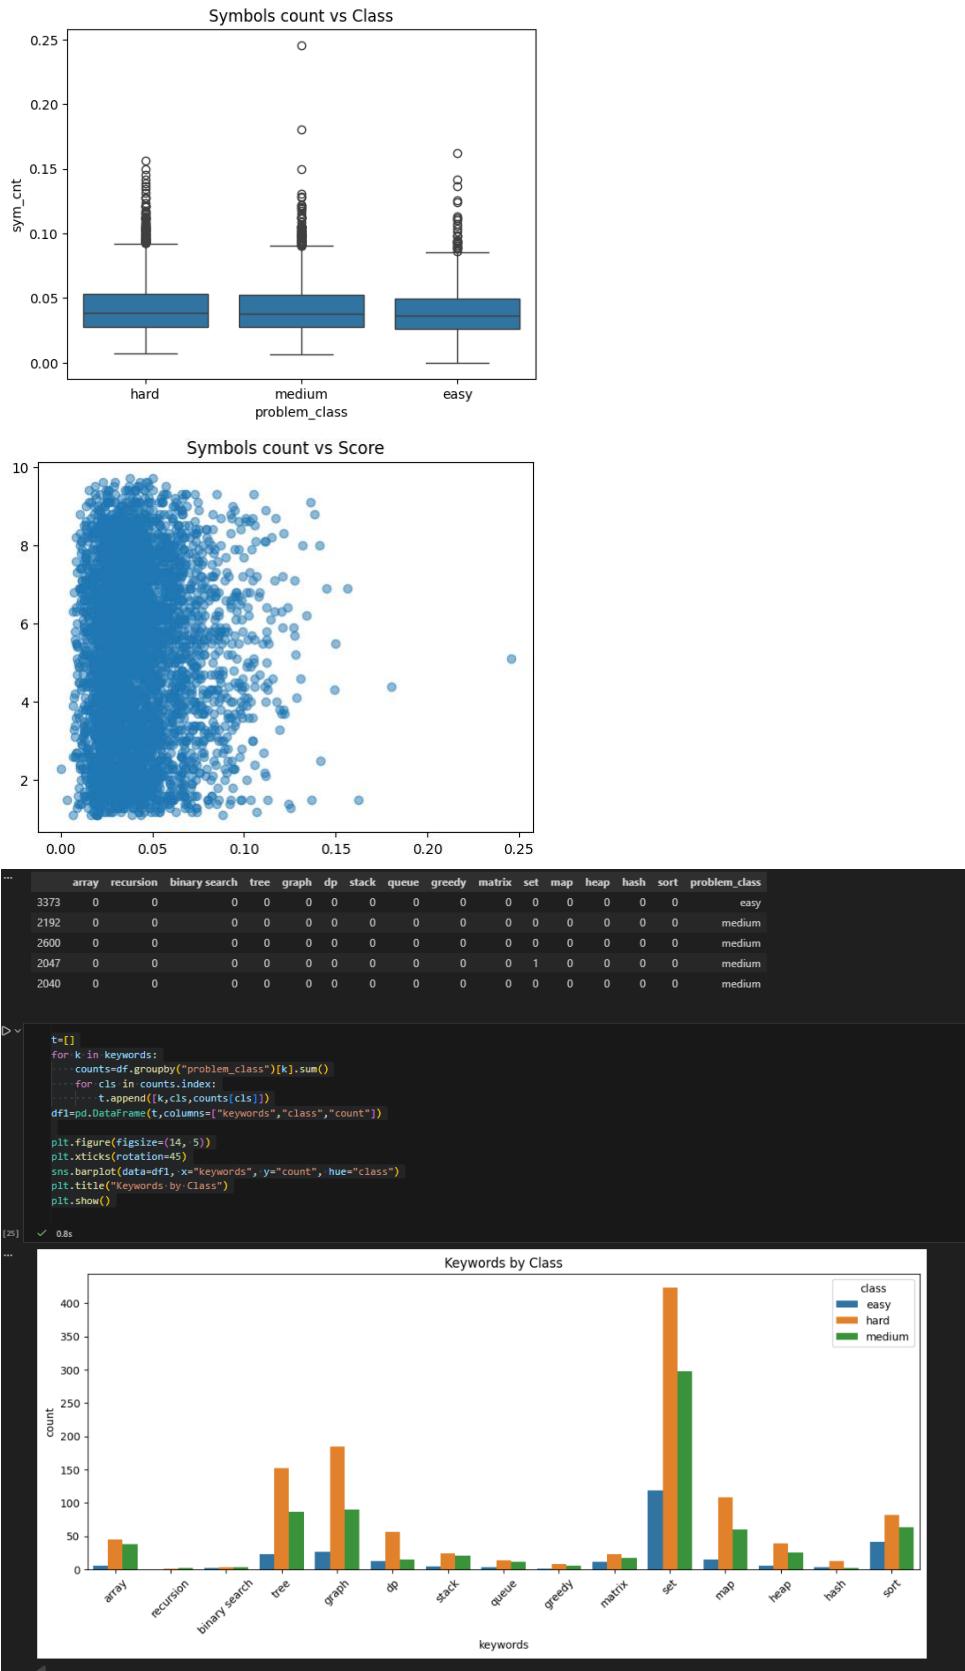
```
df["text"] = (
    df["title"].str.lower().str.strip() + " " + df["description"].str.lower().str.strip() + " " +
    df["input_description"].str.lower().str.strip() + " " + df["output_description"].str.lower().str.strip()
)
df["text_len"] = df["text"].str.len()
df["text_len"].describe()

count      4111.000000
mean     1623.911700
std      756.666678
min      10.000000
25%     1113.000000
50%     1514.000000
75%     2000.000000
max     7580.000000
Name: text_len, dtype: float64

plt.hist(df["text_len"], bins=20, edgecolor="black")
plt.title("Total text length Distribution")
plt.show()
```







```
27]      df['word_cnt'] = df['text'].str.split().str.len()
          df['word_cnt'].describe()
    ✓ 0.3s
..   count    4111.000000
     mean     273.041352
     std      125.220126
     min      1.000000
    25%    188.000000
    50%    255.000000
    75%    336.000000
    max    1226.000000
Name: word_cnt, dtype: float64

28]      df['unique_words'] = df['text'].str.split().apply(lambda x: len(set(x)))
          df['unique_words'].describe()
    ✓ 0.5s
..   count    4111.000000
     mean     143.362442
     std      51.788123
     min      1.000000
    25%    108.000000
    50%    138.000000
    75%    173.000000
    max    418.000000
Name: unique_words, dtype: float64
```

- Models Trained

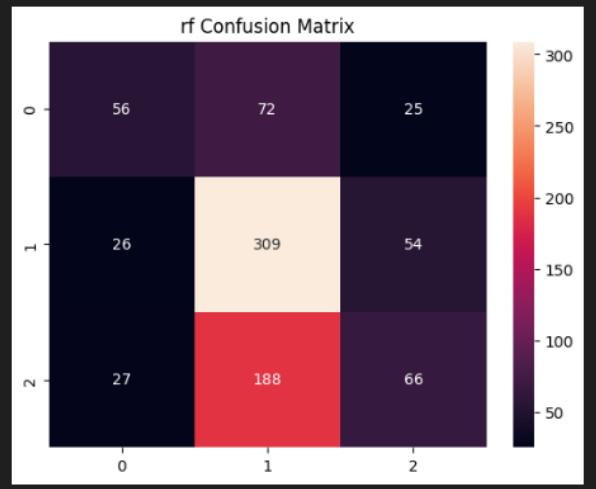
Classification Models (Predict Easy/Medium/Hard)

1. I trained 3 different models and compared accuracy:
 - Logistic Regression- 47%
 - Linear SVM -41.8%
 - Random Forest – 52.4%
2. I Used `class_weight='balanced'` to handle imbalance and Achieved 52.4% accuracy on test set (823 problems)
3. Classification Report (Random Forest):

```

rf: Accuracy = 0.524
      precision    recall   f1-score   support
      easy       0.51     0.37     0.43     153
      hard       0.54     0.79     0.65     389
      medium     0.46     0.23     0.31     281
      accuracy    0.50     0.47     0.46     823
      macro avg   0.50     0.47     0.46     823
      weighted avg 0.51     0.52     0.49     823

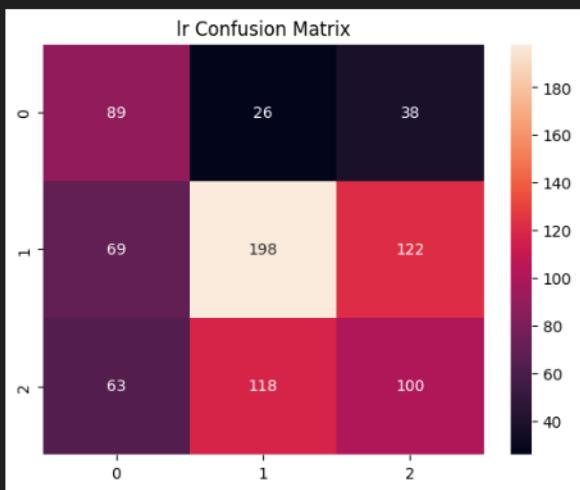
```



```

lr: Accuracy = 0.470
      precision    recall   f1-score   support
      easy       0.40     0.58     0.48     153
      hard       0.58     0.51     0.54     389
      medium     0.38     0.36     0.37     281
      accuracy    0.47     0.47     0.47     823
      macro avg   0.46     0.48     0.46     823
      weighted avg 0.48     0.47     0.47     823

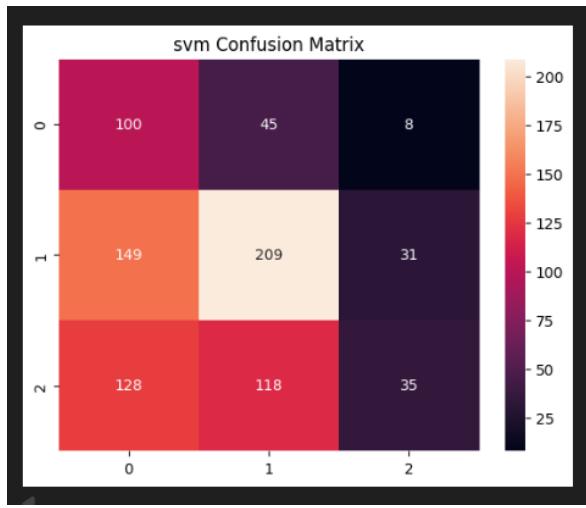
```



```

svm: Accuracy = 0.418
      precision    recall   f1-score   support
      easy       0.27     0.65     0.38     153
      hard       0.56     0.54     0.55     389
      medium     0.47     0.12     0.20     281
      accuracy    0.42     0.42     0.42     823
      macro avg   0.43     0.44     0.37     823
      weighted avg 0.48     0.42     0.40     823

```



Regression Model (Predict Score 1-10)

- I trained 3 different models and compared accuracy:

```

Linear Regression
MAE : 2.298
RMSE: 2.854
R2  : -0.68

Ridge Regression
MAE : 1.773
RMSE: 2.116
R2  : 0.077

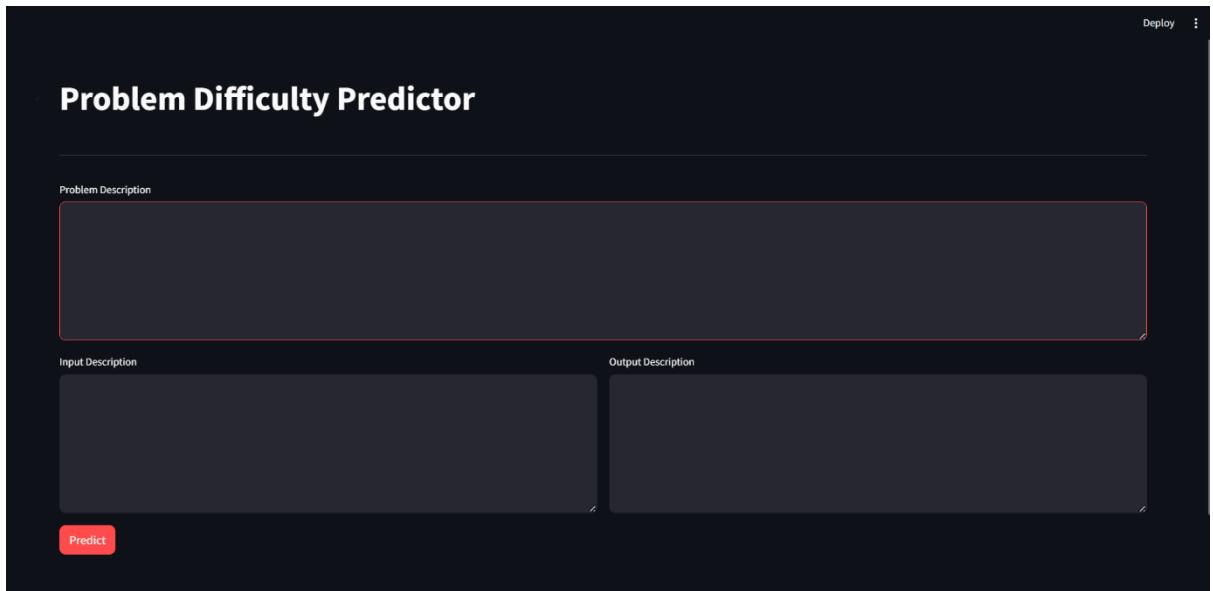
Gradient Boosting
MAE : 0.744
RMSE: 0.921
R2  : 0.825

Model Comparison:
      Model      MAE      RMSE      R2
0  Linear Regression  2.298462  2.854069 -0.679718
1  Ridge Regression   1.773196  2.115859  0.076832
2  Gradient Boosting  0.743778  0.920731  0.825187

```

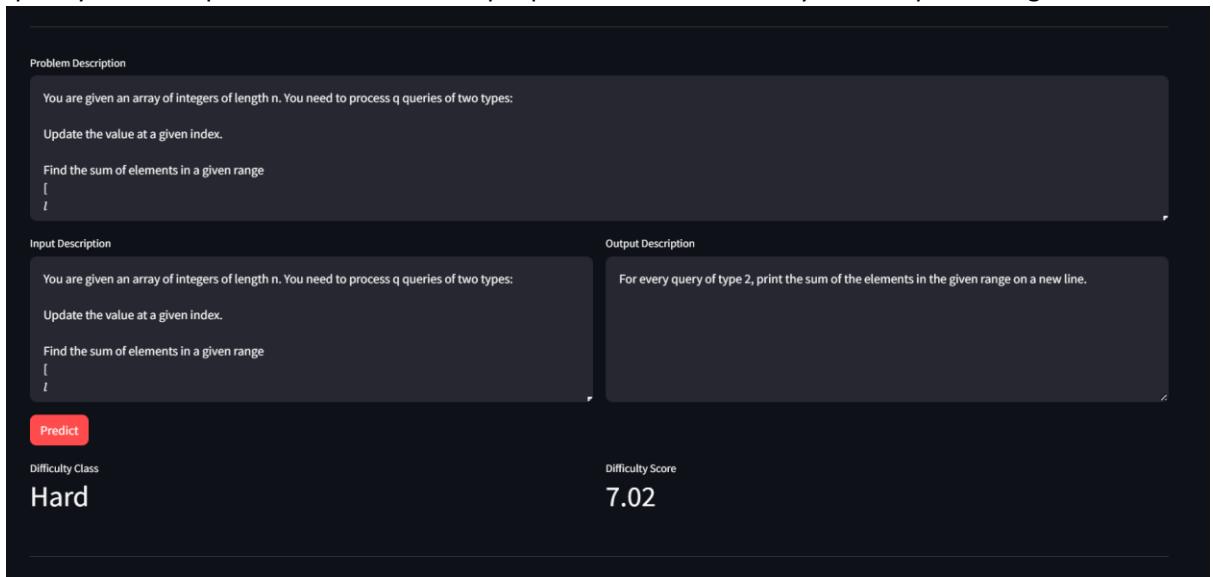
- Web interface (streamlit):

- I built a simple web application where users first paste three text fields: the problem description, input description, and output description. After entering the text, the user clicks the “Predict” button, and the system immediately displays the results, including the difficulty class (easy, medium, or hard) and a numerical difficulty score on a scale from 1.0 to 10.0.



2. How it works:

Internally, the system follows a streamlined prediction pipeline. First, all three user-provided inputs are combined into a single unified text block, ensuring that the complete problem statement is captured. This text is then processed using the same TF-IDF transformation that was applied during model training, guaranteeing consistency between training and inference. The transformed features are passed to the pre-trained machine learning models, where one model performs classification to predict whether the problem is Easy, Medium, or Hard, while the other performs regression to estimate a numerical difficulty score. Finally, the system displays the predictions instantly through the web interface, allowing users to quickly test new problems and view sample predictions without any manual processing.



● Conclusion

1. What I Achieve:

I built a classification model that predicts whether a programming problem is Easy, Medium, or Hard, achieving 51% accuracy. I also developed a regression model to predict numerical difficulty scores. In addition, I created a working Streamlit web interface, and the system can make predictions for any new problem text provided by the user.

2. Key Findings:

I found that Random Forest performed best among the tested models, achieving the highest classification accuracy. The model is particularly strong at identifying hard problems, with a recall of around 90%. I observed that text length and TF-IDF features are strong indicators of problem difficulty. However, the imbalanced dataset negatively affects predictions for easy and medium problems.

3. Why the Accuracy Is Only 52%?

This task involves three-class classification, which is inherently more challenging than binary classification. The dataset is highly imbalanced, with significantly more hard problems than easy ones. Additionally, the model relies only on text-based features, which limits how well difficulty can be captured.

4. Limitations:

The system does not use explicit problem tags or topics (such as recursion or graphs) as labels. It also depends entirely on text features and does not include other problem-level metrics. Furthermore, the imbalanced training data reduces performance for underrepresented classes.

5. Future Improvements:

I can improve the system by using pre-trained word embeddings such as Word2Vec or BERT and by collecting more data for easy and medium problems. Trying deep learning models like LSTMs or Transformers could help capture deeper text patterns. Other improvements include adding problem topic tags as features, fine-tuning hyperparameters, and combining predictions from multiple models.

6. Real-World Use:

I can use this system to recommend programming problems to students based on their skill level, help problem setters estimate difficulty, and support the creation of personalized learning paths. It can also speed up problem curation on competitive programming platforms by reducing reliance on manual difficulty assessment.
